

SLASH2

A Filesystem for Widely Distributed Systems

Pittsburgh Supercomputing Center
Advanced Systems Group



*Paul Nowoczynski
Jared Yanovich
Zhihui Zhang*

<http://quipu.psc.teragrid.org/slash2>

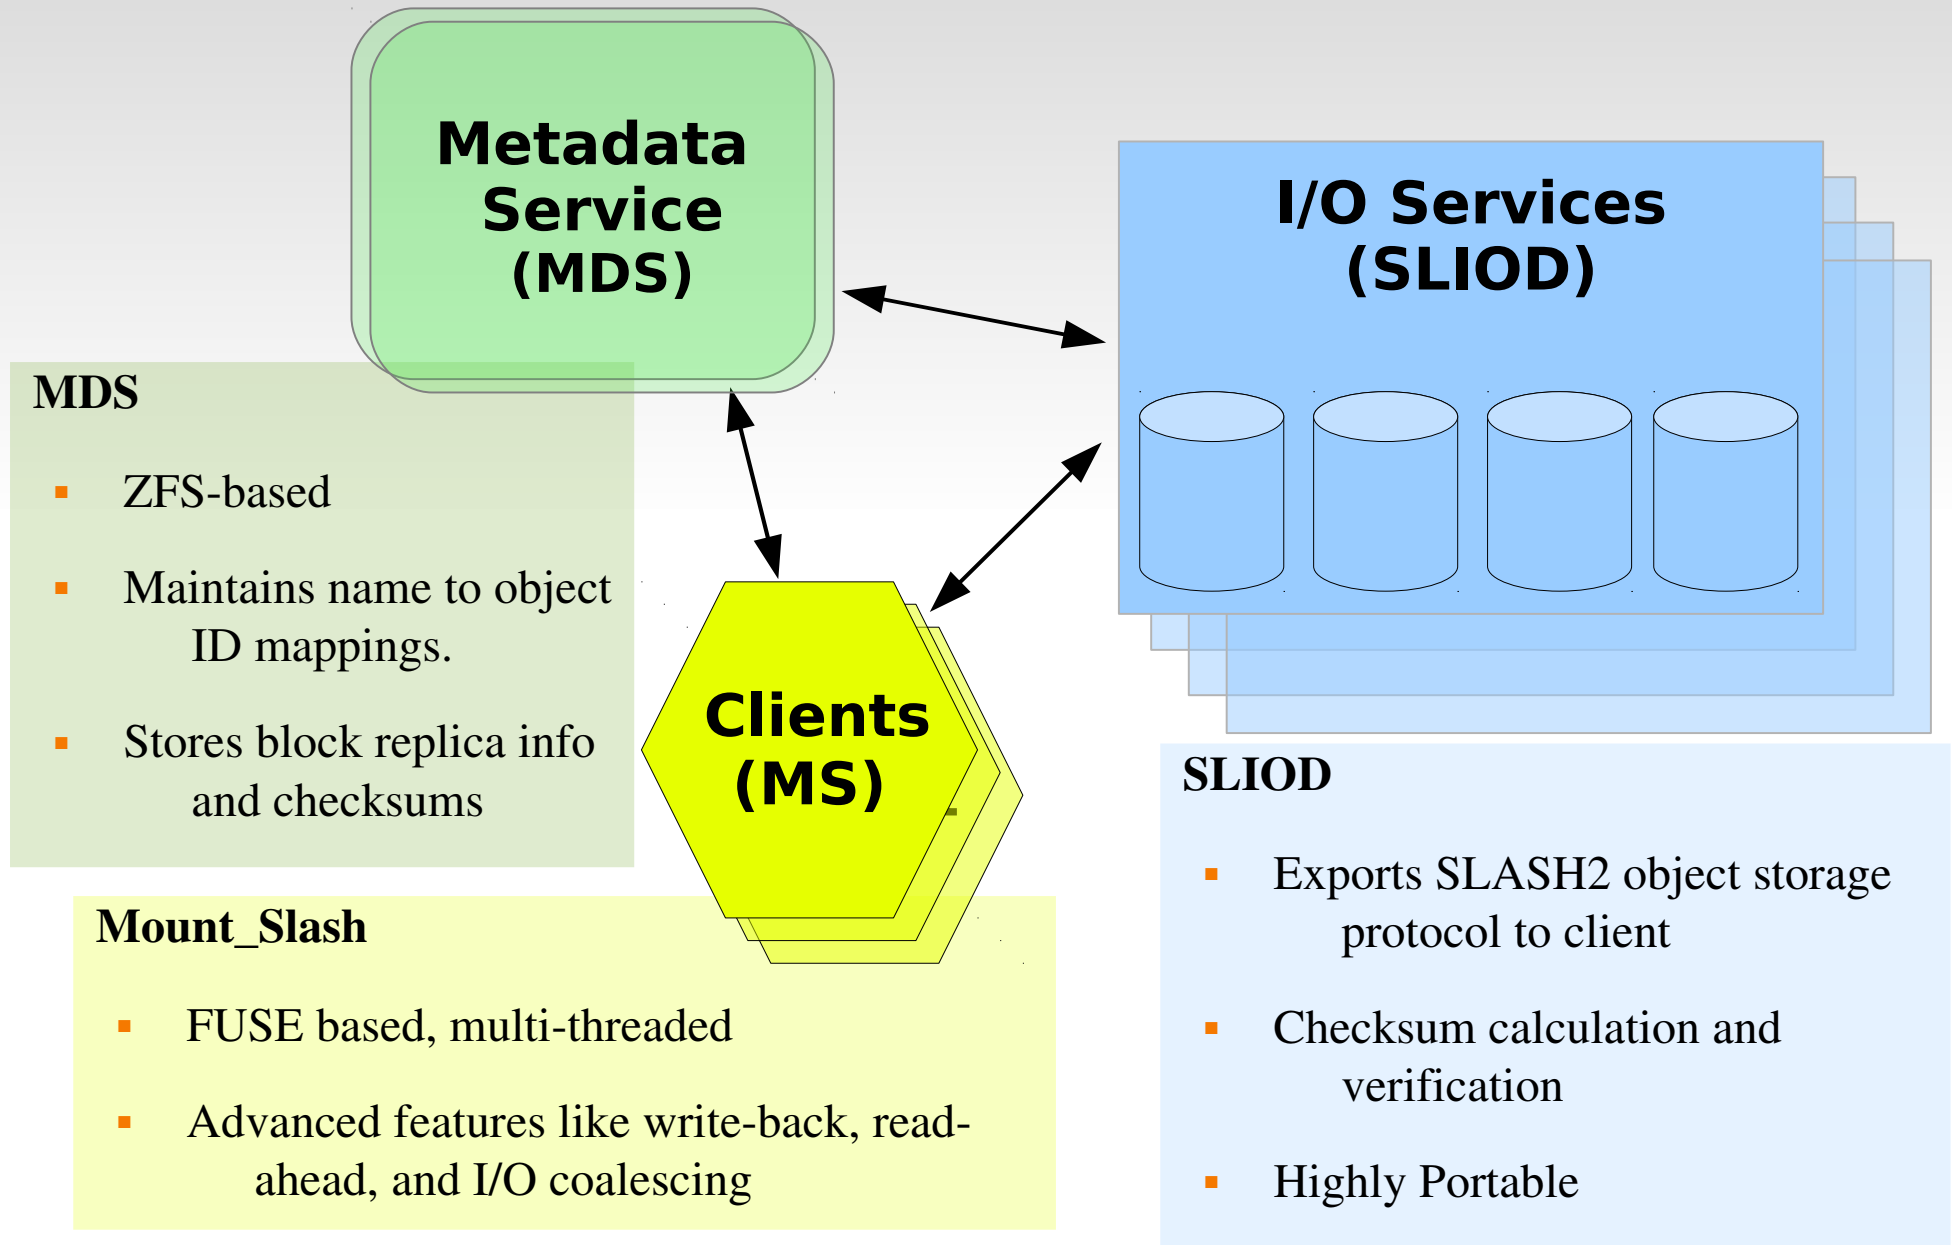
SLASH2 High-Level Architecture

- Distributed / Parallel Filesystem
 - ..with focus on wide-area
- Similar to Lustre in terms of core architecture
 - Object-based
 - Portals RPC / LNet
- Runs in user space
 - I/O services are highly portable
 - Clients utilize FUSE

Motivation behind this work

- Create filesystem which can logically bind independent storage systems
- Provide a common wide-area filesystem
 - Span institutions
 - Integrate with a variety of existing storage systems
- Alternatives solutions: Lustre-WAN, iRODS, GPFS, Panache
- Issues with existing solutions
 - too intrusive, vendor specific, or don't allow for POSIX I/O, don't support mult-resident data

SLASH2 Components



Features of Interest

- Multi-resident data
 - System managed replication
 - Supports parallelism - even for a single file
- Block-level checksumming
 - Stored and maintained on the metadata server
- Location aware data retrieval
 - Client may chose source based on proximity or other metric

Features for WAN

Filesystem protocols designed to decrease or minimize communications between clients and servers

- Create-on-write backing objects
- Readdir+
 - Minimizes attribute fetch requests for operations like: 'ls -al', 'find'..
- File size stored on metadata server
- Asynchronous garbage collection
 - Invalidated file replicas
 - Unlinked file objects
 - Fully truncated file objects

Portable I/O Service

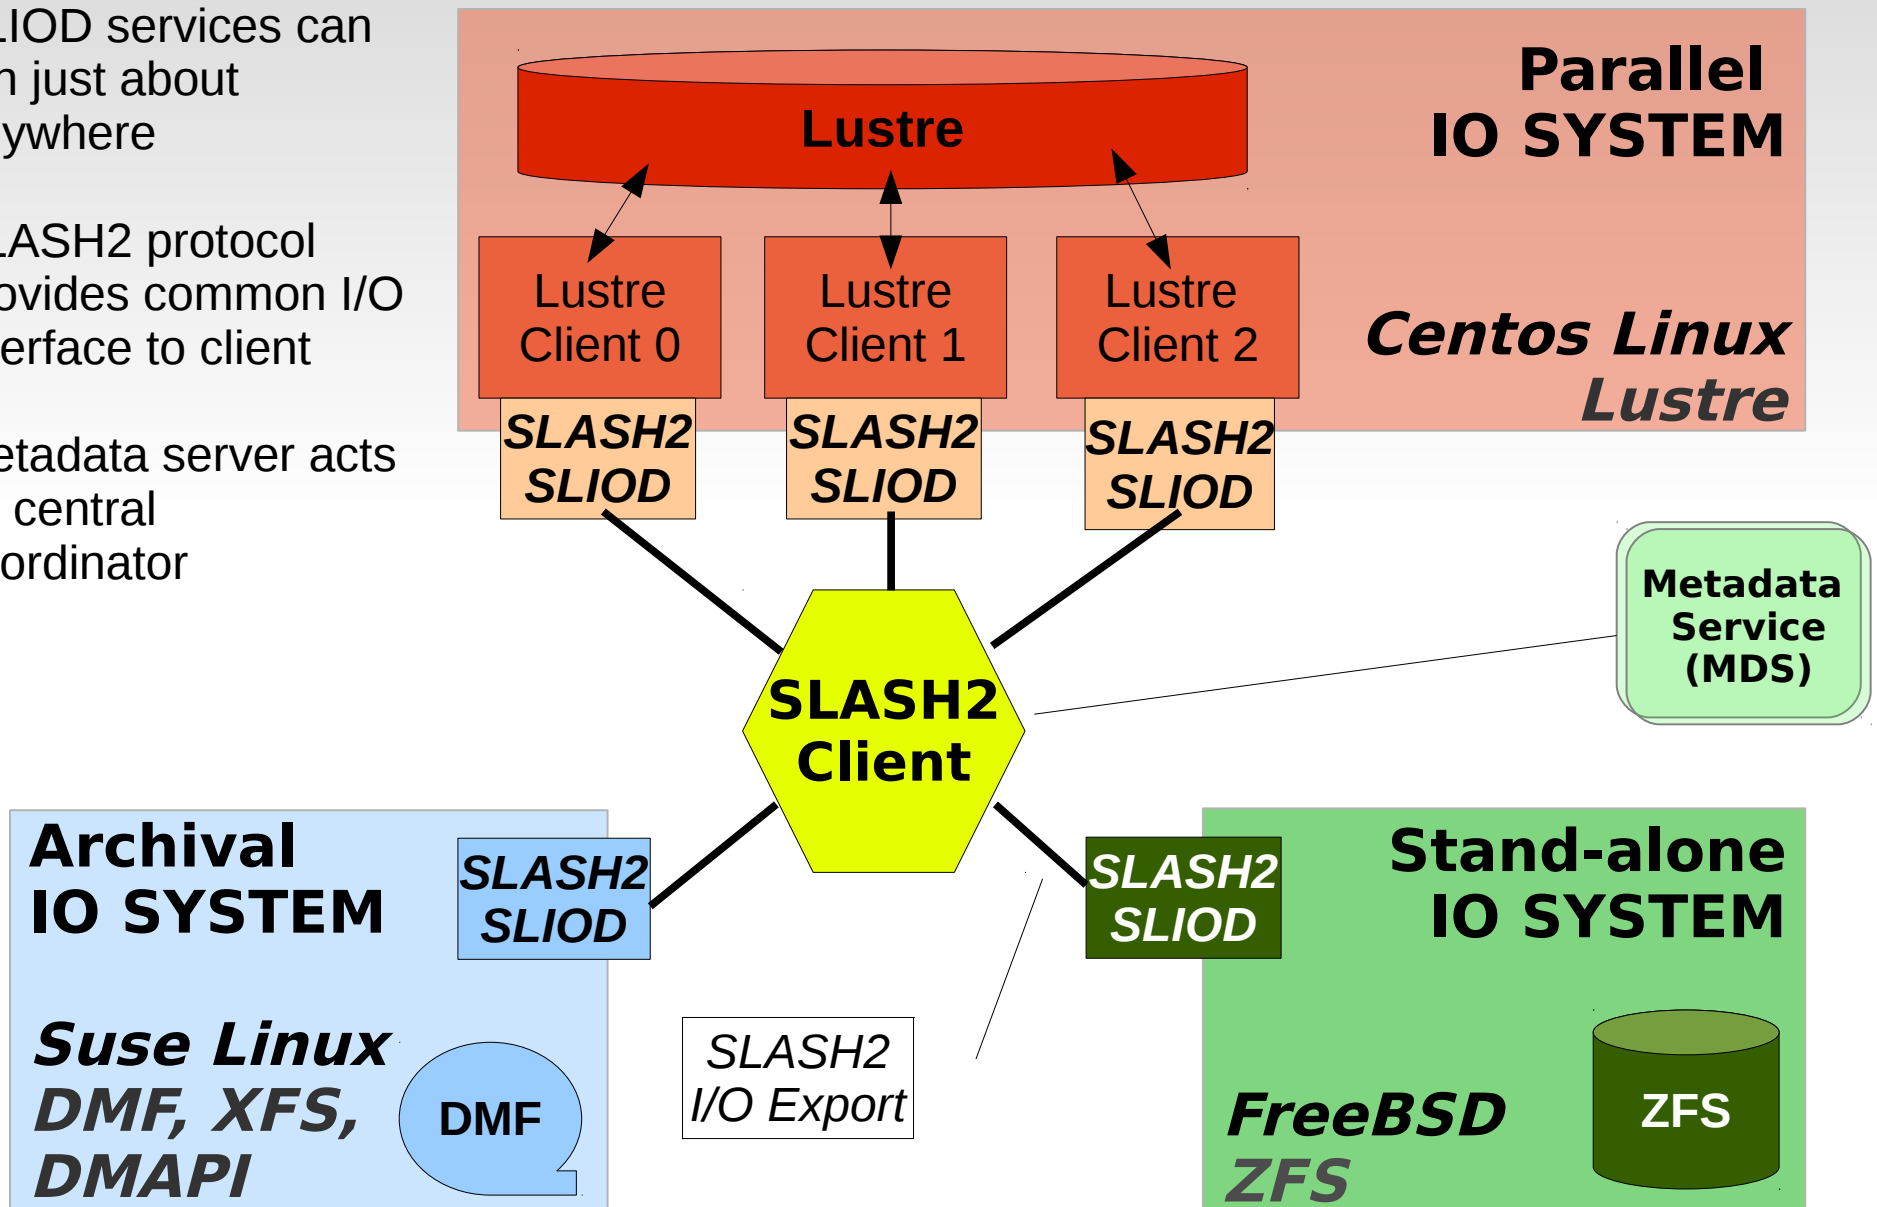
- Designed for use on a wide variety of systems
- Utilizes existing filesystems as storage backends for SLASH2
- Flexible system requirements
 - Presentation of a POSIX filesystem
 - TCP sockets
- No kernel modules required
 - Behaves like an application to the underlying FS
 - Open, close, create, unlink, p[write|read], etc.

Possible I/O System Configuration

SLIOD services can run just about anywhere

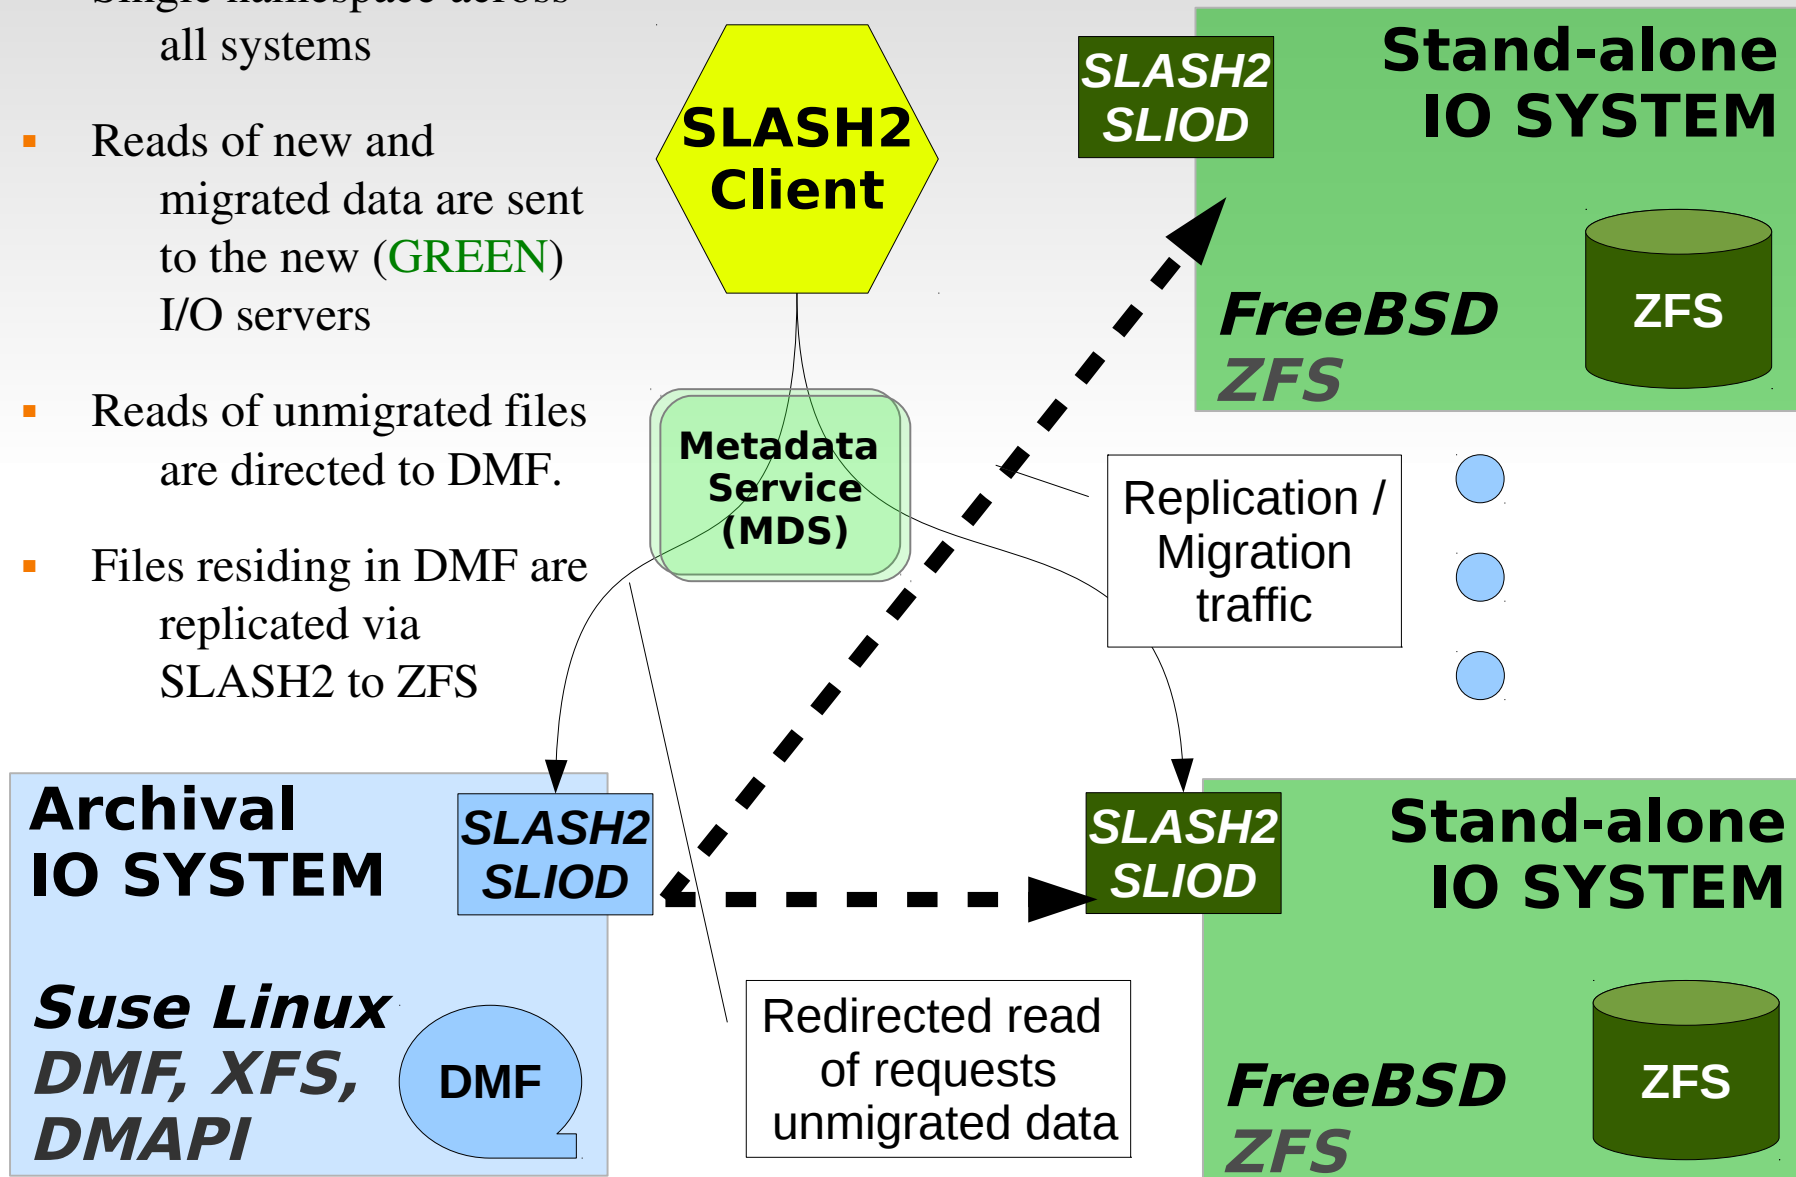
SLASH2 protocol provides common I/O interface to client

Metadata server acts as central coordinator



PSC Archival Configuration

- Single namespace across all systems
- Reads of new and migrated data are sent to the new (**GREEN**) I/O servers
- Reads of unmigrated files are directed to DMF.
- Files residing in DMF are replicated via SLASH2 to ZFS



PSC Archival Store

- Upcoming system utilizing SLASH2
- Deployment will feature core SLASH2 capabilities
- System-managed data transfer
 - Handle migration from old system to new
- Encapsulation of otherwise disparate systems
 - Logically binds a DMF / Tape-based system with a cluster of FreeBSD / ZFS servers.

SLASH2 Metadata Structures

- Unix file attributes are stored within ZFS object for performance
- SLASH2 specific data is stored as ZFS contents on the MDS
- Per file support for replication schemes and policies

SLASH2 Metadata Structures

```
/*  
* The inode structure lives at the beginning of the metafile.  
* @ino_version: compatibility.  
* @ino_flags: slash2 specific file attributes.  
* @ino_bsz: size of this objects bmap.  
* @ino_nrepls: number of replicas, if > SL_DEF_REPLICAS use inode_extras.  
* @ino_replpol: file replication policy.  
* @ino_repls: replica storage.  
* @ino_repl_nblks: support st_blocks in multi-res filesystem.  
*/
```

```
struct slash_inode_od {  
    uint16_t      ino_version;  
    uint16_t      ino_flags;  
    uint32_t      ino_bsz;  
    uint32_t      ino_nrepls;  
    uint32_t      ino_replpol;  
    sl_replica_t   ino_repls[SL_DEF_REPLICAS];  
    uint64_t      ino_repl_nblks[SL_DEF_REPLICAS];  
};
```

```
struct slash_inode_extras_od {  
    sl_replica_t   inox_repls[SL_INOX_NREPLICAS];  
    uint64_t      inox_repl_nblks[SL_INOX_NREPLICAS];  
};
```

Metadata & Multi-Residency

- SLASH2 maintains per chunk metadata
- 'Chunks' (32MiB-256MiB) are bigger than 'blocks' (4KiB-4MiB)
- The structure which describes chunks is called a *bmap*
- Residency and policy and managed per *bmap*

```
/*  
 * @bmod_crcstates: bits describing the state of each sliver.  
 * @bmod_repls: bitmap used for tracking the replication status of this bmap.  
 * @bmod_crcs: the CRC table, one 8-byte CRC per sliver.  
 * @bmod_gen: current generation number.  
 * @bmod_replpol: replication policy.  
 */  
struct bmap_ondisk {  
    uint8_t        bmod_crcstates[SLASH_CRCS_PER_BMAP];  
    uint8_t        bmod_repls[SL_REPLICA_NBYTES];  
    uint64_t       bmod_crcs[SLASH_CRCS_PER_BMAP];  
    uint32_t       bmod_gen;  
    uint32_t       bmod_replpol;  
}
```

Replication Example

- Create a file

```
(pauln@peel0:~)$ dd if=/dev/zero of=/p0_archive/pauln/big_file count=2k bs=1M
2048+0 records in
2048+0 records out
2147483648 bytes (2.1 GB) copied, 4.81828 seconds, 446 MB/s
```

- Note current residency status

- File lives on resource “archsliod@PSCARCH”

```
(pauln@peel0:msctl)$ ./msctl -r /p0_archive/pauln/big_file
file-replication-status                                     #valid  #bmap  %prog
=====
/p0_archive/pauln/big_file
new-bmap-repl-policy: one-time
  archsliod@PSCARCH
  ++++++
```

*16 bmaps resident at
I/O resource
archsliod@PSCARCH*

16 16 100%

- Issue replication request to resource “archlime@PSCARCH”

```
(pauln@peel0:msctl)$ date && ./msctl -Q archlime@PSCARCH:*/p0_archive/pauln/big_file
Wed Jul 20 02:51:56 EDT 2011
```

Replication Example (2)

■ Check status

Wed Jul 20 02:51:57 EDT 2011

file-replication-status

=====

/p0_archive/pauln/big_file
new-bmap-repl-policy: one-time
archsl iod@PSCARCH
+++++

archlime@PSCARCH
sqqqqqqqqqqqqqqqqq

*1 second elapsed
1 bmap en route,
15 enqueued*

#valid	#bmap	%prog
16	16	100%
0	16	0%

Wed Jul 20 02:52:05 EDT 2011

file-replication-status

=====

/p0_archive/pauln/big_file
new-bmap-repl-policy: one-time
archsl iod@PSCARCH
+++++

archlime@PSCARCH
+++++qqqqqq

*8 seconds elapsed
10 bmaps done
6 enqueued*

#valid	#bmap	%prog
16	16	100%
10	16	62.50%

Wed Jul 20 02:52:20 EDT 2011

file-replication-status

=====

/p0_archive/pauln/big_file
new-bmap-repl-policy: one-time
archsl iod@PSCARCH
+++++

archlime@PSCARCH
+++++

*24 seconds elapsed
all bmaps done !*

#valid	#bmap	%prog
16	16	100%
16	16	100%

More on Replication..

What just happened? System Managed Data Xfer

1. The 'msctl' command created a replication request

```
(pauln@peel0:msctl)$ date && ./msctl -Q archlime@PSCARCH:*/p0_archive/pauln/big_file  
Wed Jul 20 02:51:56 EDT 2011
```

2. This caused the MDS to statefully create a work request for 'big_file'. Once complete.. it returned 'OK' to the client
3. MDS scans the metadata for bmaps which do not meet the new replication criteria, making a work item for each.
4. Work items are scheduled to be sent from

'archslod@PSCARCH' → 'archlime@PSCARCH'

Upon receiving the work request, the SLIOD retrieves the checksum table from the MDS to verify incoming contents.

5. MDS provides some flow control and rudimentary scheduling – as work is completed, more is allocated until done

Replication - What wasn't shown..

- When using I/O systems which have multiple SLIODs, the MDS can distribute replication work across all destination SLIODs.
- Robustness
 - MDS resumes all unfinished replication work on restart
 - MDS can cope with missing or slow SLIOD endpoints
- Load balancing
 - SLIODs are given work piecemeal.
 - Parallel replications are not disproportionately affected by a single slow or oversubscribed node.

Metadata Replication

- Provide near-uniform metadata performance throughout the wide-area
 - Lots of metadata operations rely on small, serial RPCs
 - multi-millisecond lookup()'s are performance killers!
- Eventual consistency is the only way..
 - Maintaining read and write locks on directories and metadata structure will not work
 - Wide-area latency will crush performance
 - Network partitioning and spanning of administrative domains present challenges

Metadata Replication (2)

- Current plan will rely on ZFS snapshotting
- Lots of advantages
 - zfs [send|recv] does all of the heavy lifting
 - Easily administered (via standard zfs tools)
 - New metadata servers may be easily incorporated

Metadata Replication (3)

How it works..

- Each MDS has a ZFS filesystem for himself and each of his peers
- Local MDS may only modify his designated ZFS filesystem
 - Publishes modifications to his peers
 - Receives updates from his peers on behalf
 - Asymmetric performance for some updates ..BUT..
 - Can make progress on local resources in the event of network partitions or remote server failures

Metadata Replication (4)

- Eventually consistent namespace - system aims to for a 60 second update interval
 - i.e. if the system quiesced for >60 seconds, all MDSs should reach synchronization
 - This is goal.. may not be realistic
- More details will be made available on the web..

Questions?

<http://quipu.psc.teragrid.org/slash2>